
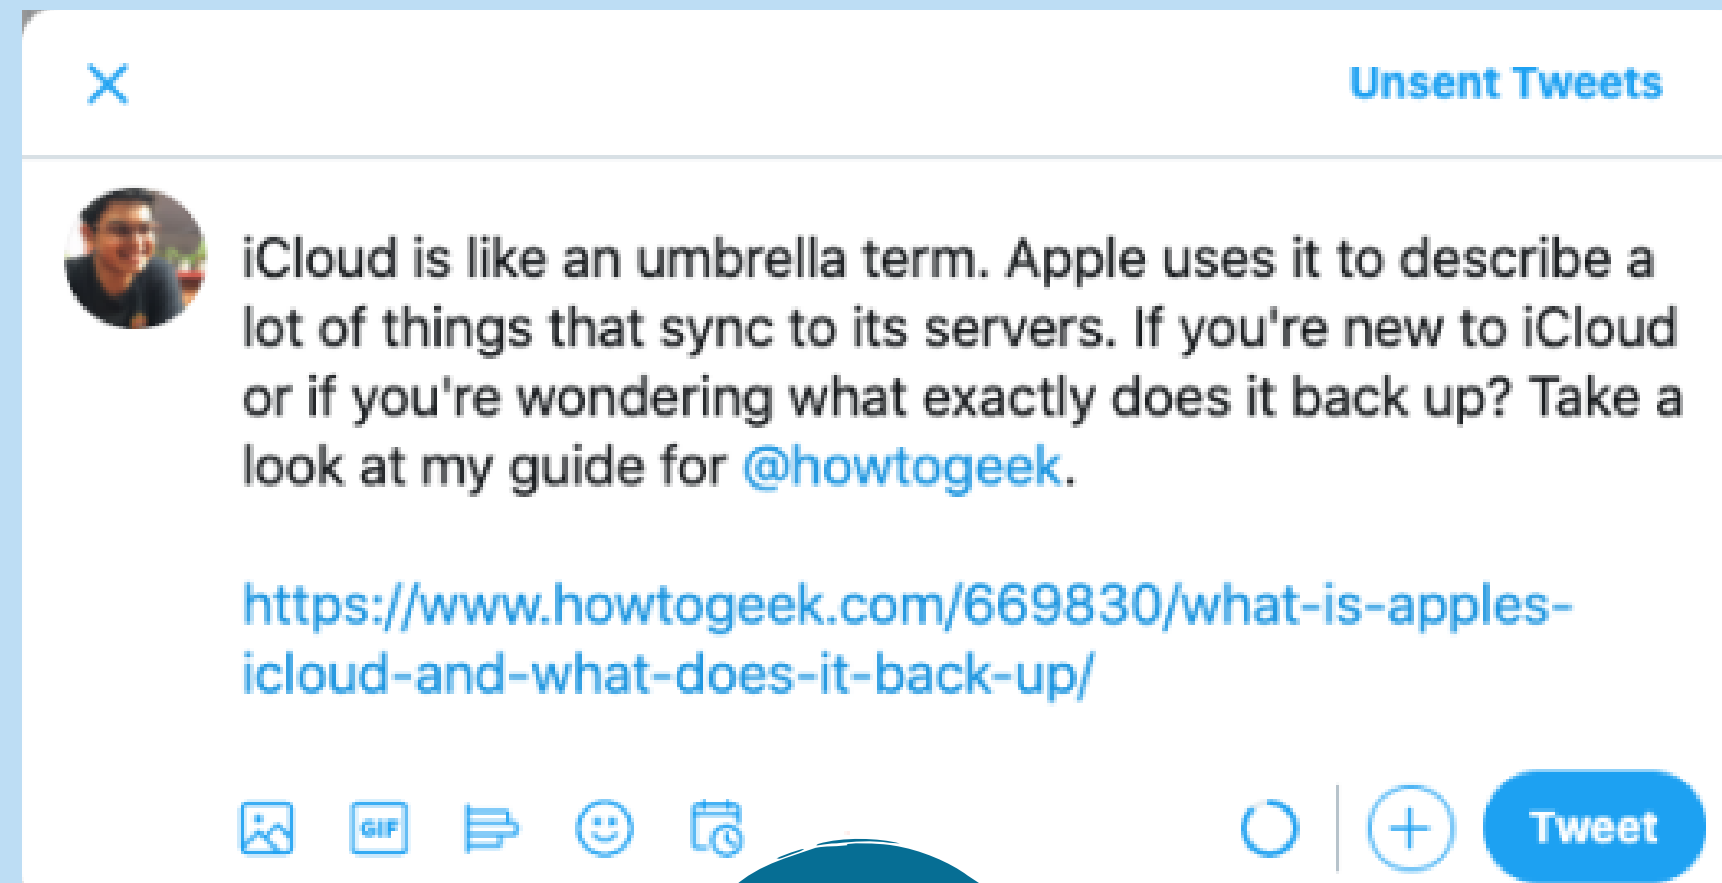


# Disseminating fake information



# Project Goal



# Steps

- 1 Loading the data
- 2 Retrieving information
- 3 Pre-processing
- 4 Feature Importance
- 5 ML Models
- 6 Choosing the Classifier

# Loading the data

sklearn  
+  
pandas } libraries  
initially used

This is where we are  
retrieving our  
information from

85% - 15% split

Hidden



The diagram illustrates a data split. A light blue arrow points from the text 'This is where we are retrieving our information from' to the text '85% - 15% split'. In this split, the '15%' is circled in light blue. A curved light blue arrow points from the word 'Hidden' to the circled '15%', indicating that the 15% of the data is the hidden set.

# Retrieving information

## Most important words, hashtags, and mentions

### 1. Separate the tweets

```
true_tweets = tweets.loc[tweets['questionable_domain'] == 0]  
fake_tweets = tweets.loc[tweets['questionable_domain'] == 1]
```

### 2. Retrieve all the words, hashtags, and mentions from each group



2.1. Make a top n most used words, hashtags or mentions for each group, based on its frequency among the tweets of each group

### 3. Retrieve all links from each group

3.1. For each tweet-like link, convert it to visible domain

3.2. For each visible domain, retrieve the root domain

# Retrieving information



# Retrieving information

## Hashtags Word Clouds



"Good" Hashtags



"Bad" Hashtags

# Pre-processing

## 1 Define auxiliary functions

### To process text:

- count words
- count words in CAPS
- retrieve hashtags
- retrieve root domains
- retrieve mentions
- retrieve the most common words
- evaluate emotion levels
- sentiment analysis

libraries used  
for pre-processing { text2emotion  
nltk  
re



# Pre-processing



# Pre-processing

**Sentiment Analysis** – few steps to get there:

1. Lower case every word
2. Remove punctuation
3. Remove Stop-Words
4. Lemmatization



**Score**  $\in [-1, 1]$

# Pre-processing

## ② Define a Pre-processing function

Using every auxiliary function previously defined we mounted some sort of a **pipeline for pre-processing** any data frame that had the same structure.

It goes from **removing unused features** like the identifier and duplicated features in content, to **performing some NLP Techniques** in order to perform both **Emotion Extraction** and **Sentiment Analysis**.

This is where feature engineering is done: **feature extraction** and **feature selection**!

# Feature Importance

We used a classifier named ExtraTreesClassifier from **sklearn** library to perform a feature importance measurement.

We did just like we were to fit a model, but we used the **feature\_importances\_** method that comes with this classifier.

```
X = data.loc[:, data.columns != 'questionable_domain']

y = data['questionable_domain']

# feature extraction
model = ExtraTreesClassifier(n_estimators=10)
model.fit(X, y)

importance_list = list(model.feature_importances_)
```

# Feature Importance

```
a = features_names[1:]
b = importance_list

c = zip(a,b)
c = list(c)

d = sorted(c, key = lambda x: x[1])

to_delete = []
for i in range(len(d)):
    if d[i][1] <= 0.01:
        to_delete.append(d[i][0])

tweets = tweets.drop(to_delete, axis=1)
```

We zipped the features' names and their score of feature importance so we could know which features to remove, in order to perform **dimensionality reduction**.


We, then, **sorted** the list and **filtered** it to have only features that had an importance score **greater than 1%**.

A list named **to\_delete** was created to facilitate the drop features step.

# ML Models

Naive Bayes

AdaBoost



Random Forrest

Voting Classifier

# ML Models

**Before fitting the models:**

## Imbalanced Data



```
#check out how many fake and true cases we have:  
print(tweets['questionable_domain'].value_counts())
```

```
0    12720  
1     2537
```

random  
under-sampling



majority class


random  
over-sampling



minority class



True  
50%



Fake  
50%

# ML Models

## Evaluating the models:

# 1 X 10 CROSS VALIDATION

## Naive Bayes

accuracy: 0.6207621550591327  
precision: 0.6069025654781756  
recall: 0.7543721508495648  
f1: 0.6611756459412084

## AdaBoost

accuracy: 0.8169513797634691  
precision: 0.8730174228175306  
recall: 0.7419242989363172  
f1: 0.8020081100018412

## Random Forrest



accuracy: 0.9195795006570302  
precision: 0.951277281667301  
recall: 0.8856858682138418  
f1: 0.9166526091862528

## Voting Classifier

accuracy: 0.8629434954007884  
precision: 0.8943853852335609  
recall: 0.8263157894736841  
f1: 0.8595269458416139



# Choosing the Classifier

The classifier with the best performance, according to the chosen evaluation method is **Random Forrest**.

Now, it's time to fit/train our chosen model and use it in our **hidden** data frame (the 15% split we did at the beginning).

This step is done with an already balanced train data frame, so our model can recognize as many true cases as fake cases.

# Choosing the Classifier

One other "small" detail that we cannot forget about it:

The hidden data frame is not yet pre-processed at this stage, and this is why it was so convenient to create a pre-processing function.

At this point, we just need to pre-process it and we're all set to go!

```
hidden_after_preProcessing = preProcessing(hidden)
```

# Choosing the Classifier

After the pre-processing on hidden is done, we separate the target variable from the rest, and we let our model predict the **unseen** tweets.

The results are as follow:

```
RandomForestClassifier() accuracy: 0.9220200519866321  
RandomForestClassifier() precision: 0.7962529274004684  
RandomForestClassifier() recall: 0.734341252699784  
RandomForestClassifier() f1: 0.7640449438202248
```